



---

## **Understanding Stock Price Dynamics & Tailoring Investment Approaches of Tata Consultancy Services Limited**

---



MSc FinTech and Business Analytics

Module: Artificial Intelligence and Machine Learning

Module Code: FNCE043W

Module Leader: Dr. Yumei Yao

Student Name: Sarika Patel

Student ID: W2019695

Word Count: 2183 words (Excluding cover page, Index, References, Appendix)

## TABLE OF CONTENT

### CONTENTS

|  |    |
|--|----|
| Abstract .....   | 2  |
| 1.Introduction.....  | 2  |
| 2.Data Collection and processing .....                     | 2  |
| a. Historical Data.....                                    | 2  |
| b. Feature Engineering .....                               | 2  |
| 3.Exploratory Data Analysis of Features .....              | 4  |
| 4.Machine Learning Classification .....                    | 8  |
| a.Logistic Regression .....                                | 8  |
| b. Extra Trees .....                                       | 9  |
| 5.Cross Validation Accuracy Analysis .....                 | 9  |
| a. cross validation for Logistic Regression .....          | 9  |
| b. cross validation for Extra Trees .....                  | 10 |
| 6.Evaluation of Classifiers Results .....                  | 10 |
| 7.Prediction of Price using X_ Test Data: Extra Tree ..... | 11 |
| a. Classification report.....                              | 11 |
| b. Confusion Matrix .....                                  | 12 |
| c.ROC Curve .....  | 12 |
| 8. Market and Strategy Returns .....                       | 13 |
| 9.Cummulative Returns Analysis .....                       | 13 |
| 10.Interpretation and Discussion .....                     | 14 |
| 11.Conclusion .....  | 15 |
| 12.References.....   | 15 |
| 13.Appendices.....   | 16 |

## ABSTRACT

This study examines how Tata Consultancy Services (TCS) stock price changes were predicted using machine learning approaches, namely Logistic Regression and Extra Trees classification algorithms, between January 2014 and December 2023. The study finds patterns and trends in historical stock data using in-depth Exploratory Data Analysis (EDA), which informs the forecast models. The research evaluates the models' performance using cross-validation, calculating accuracy, precision, recall, and F1-scores. It also compares the cumulative returns of conventional market strategies with those based on machine learning forecasts. This brief study casts doubt on machine learning's ability to anticipate stock price increases properly and offers a critical perspective on the technology's predictive capabilities in the financial industry. It also adds to the larger conversation about the application of machine learning.

## 1.INTRODUCTION

Tata Consultancy Services Limited (TCS) is an Indian multinational information technology (IT) services and consulting company headquartered in Mumbai. It functions in 150 sites around 46 countries and is a member of the Tata Group. It was said in September 2023 that TCS employed more than 616,000 people globally. TCS is the most valuable IT service brand globally, the second-largest Indian firm by market value, and the leading Big Tech (India) company. The company has generated consolidated revenues of US \$27.9 billion in the year ended March 31, 2023, and is listed on the BSE and the NSE in India. (Tata Company, 2024)

## 2.DATA COLLECTION AND PROCESSING

### A. HISTORICAL DATA

Collection of Tata Consultancy Services (TCS) daily stock data spanning from January 2014 to December 2023 was conducted utilizing Yahoo Finance. The dataset comprises 2444 rows, encompassing essential metrics such as open, high, low, and close prices, as well as adjusted close prices and stock volume. Data that was incomplete or missing with rows was eliminated.

### B. FEATURE ENGINEERING

Generated featured from the processed data include:

**i.** The daily high-low ('H-L') and open-close ('O-C') differentials represent intraday price volatility and changes in stock prices throughout the trading day.

Formula: H-L= High Price- Low Price

O-C= Closing Price- Opening Price

By monitoring H-L and O-C differences, market participants gain valuable insights into market dynamics, price volatility, and potential trading opportunities, enhancing their decision-making processes in the financial markets.

**ii.** Moving averages ('3d MA', '10d MA', and '30d MA') show how the closing price changes over short, medium, and longer periods.

Formula: n-Day MA = Sum of closing prices for the last n days / n

By analysing these moving averages, market participants gain insights into short and medium-term price dynamics.

**iii.** The closing price's standard deviation ('Std\_dev') measures daily price volatility over a 5-day rolling window.

Formula:

$$\text{Std\_dev}_t = \sqrt{\frac{\sum_{i=1}^t (X_i - \bar{X})^2}{t}}$$

Where:

- $X_i$  is each closing price at time  $i$ .
- $\bar{X}$  is the mean closing price up to time  $t$ .
- The summation is performed over all closing prices up to time  $t$ .

This rolling window approach allows us to calculate the standard deviation dynamically over successive 5-day periods, capturing short-term price volatility in the financial markets.

iv. 'Price\_Rise' - binary column which predicts price increases based on historical data.

$$\text{Price\_Rise}_t = \begin{cases} 1 & \text{if } \text{Close}_{t+1} > \text{Close}_t \\ 0 & \text{otherwise} \end{cases}$$

Where:

- $\text{Close}_{t+1}$  is the closing price at the next time step.
- $\text{Close}_t$  is the closing price at the current time step.
- $\text{Price\_Rise}_t$  is the binary column indicating whether the price rises (1) or not (0).


The 'Price\_Rise' column is a binary indicator used to predict price movements based on historical data. Typically, it assigns a value of 1 when the price is predicted to rise and a value of 0 when the price is predicted to either stay the same or fall.

By removing rows with missing values and choosing the most pertinent TCS fields for analysis, the dataset was further refined.

### 3.EXPLORATORY DATA ANALYSIS OF FEATURES

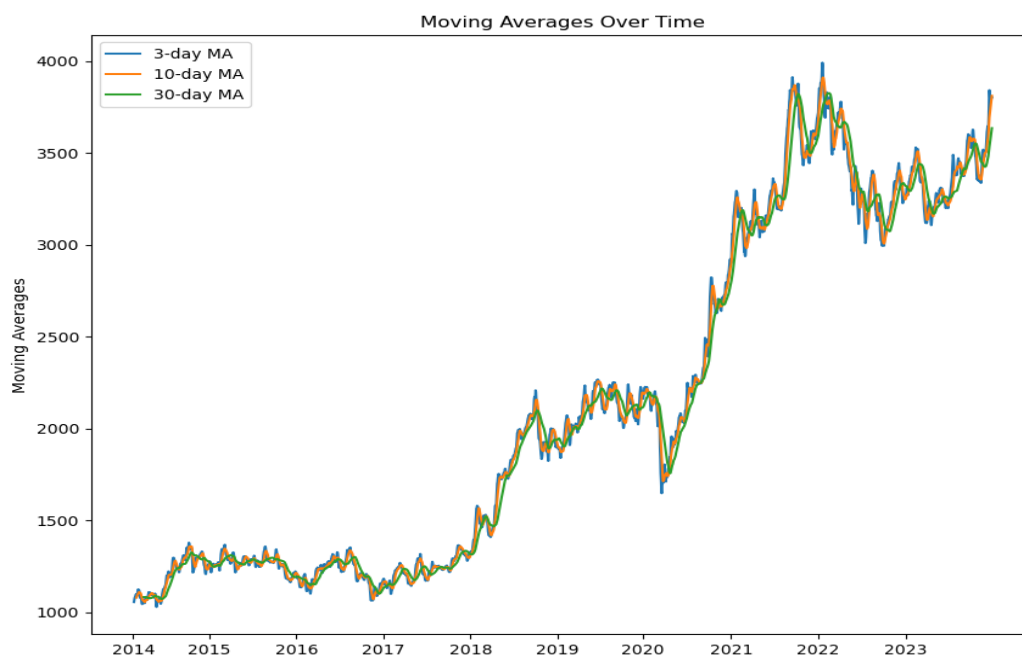
Before using advance statistical approaches, exploratory data analysis, or EDA, is a vital first stage in data analysis that aims to comprehend the key traits and patterns included within a dataset. In EDA, analysts usually look at the variables' distribution, central tendency, and dispersion, spot any outliers or missing data, and use graphs or charts to show the correlations between the variables. (IBM, 2024) EDA can provide insights into the data's structure, any correlations between variables, and any possible problems that would need to be fixed before additional modelling or analysis.

Figure 1: Closing price and volume Relationship.



The chart depicts TCS stock's progress, showing an upward trajectory in the closing price against trading volume peaks. These peaks might reflect key events but do not consistently correlate with the stock's price, suggesting a complex interplay between volume and price dynamics. (Fig.1)

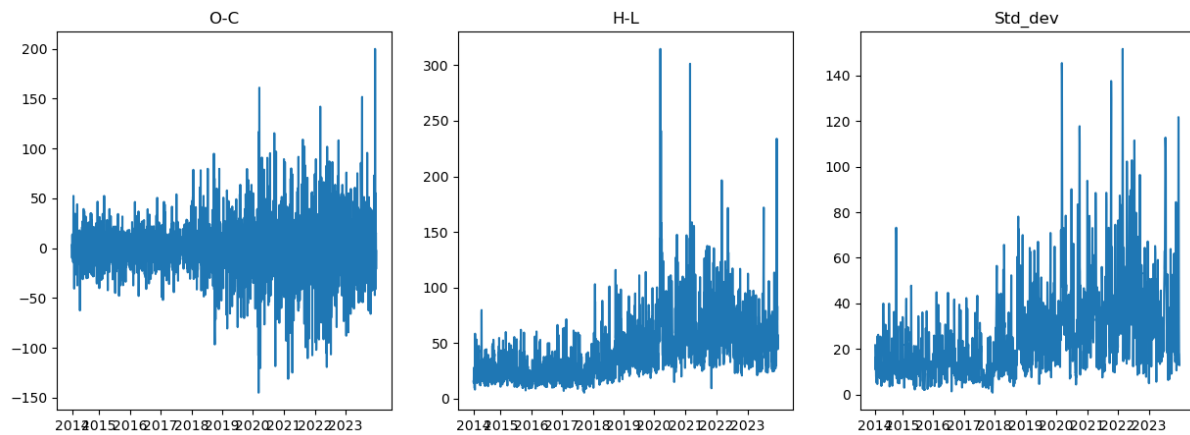
Figure 2: Plot for 3day, 10day, 30day MA



The moving averages such as '3-day MA', '10-day MA', and '30-day MA' help reduce short-term fluctuations in data. These moving averages exhibit comparable trends over time (Fig.1)

and display similar distribution patterns (Fig.4). As the rolling window widens, the resulting trend line appears smoother, coinciding with a general upward trend in closing prices, with values frequently exceeding INR 3500 across the observed timeframe.

Figure:3. O-C, H-L and std\_deviation



The 'H-L' and 'O-C' functions provide data on price fluctuations and market sentiment. The market condition in "H-L," also referred to as the daily price range, indicates that it normally rises with time, and the O-C price difference is larger as time goes on (Fig.3). The daily standard deviation moves in the same way as the "H-L" characteristic, which indicates that a larger "Std\_dev" corresponds to a wider range of possible price movement and a higher degree of risk.

Figure 4: 3day,10day,30days MA histogram and density





Figure 5: Price Rise & Fall Count Plot in relation to 'Close' and 'Std\_dev'



The price rise and fall binary column (derived from our dataset's 'Price\_Rise' feature; see Fig. 5) exhibits a negligible difference in the counts of upward and downward movements. The closing price and standard deviation predictions are generally somewhat positive oriented, showing significant consistency in some regions where the data points are densely clustered.

Figure 6: Descriptive Statistics

|      | count       | Open        | High        | Low         | Close       | Adj Close   | % |
|------|-------------|-------------|-------------|-------------|-------------|-------------|---|
| mean | 2415.000000 | 2415.000000 | 2415.000000 | 2415.000000 | 2415.000000 | 2415.000000 |   |
| std  | 2126.355000 | 2126.355000 | 2126.355000 | 2126.355000 | 2126.355000 | 2126.355000 |   |
| min  | 916.531878  | 916.531878  | 916.531878  | 916.531878  | 916.531878  | 916.531878  |   |
| max  | 1057.000000 | 1057.000000 | 1057.000000 | 1057.000000 | 1057.000000 | 1057.000000 |   |
| 50%  | 1262.712524 | 1262.712524 | 1262.712524 | 1262.712524 | 1262.712524 | 1262.712524 |   |
| 75%  | 1977.000000 | 1977.000000 | 1977.000000 | 1977.000000 | 1977.000000 | 1977.000000 |   |
| max  | 4045.500000 | 4045.500000 | 4045.500000 | 4045.500000 | 4045.500000 | 4045.500000 |   |

|      | count        | Volume      | H-L         | O-C         | 3d MA       | 10d MA      | % |
|------|--------------|-------------|-------------|-------------|-------------|-------------|---|
| mean | 2.415000e+03 | 2415.000000 | 2415.000000 | 2415.000000 | 2415.000000 | 2415.000000 |   |
| std  | 1.682678e+05 | 42.233925   | 1.399579    | 212.824875  | 212.824875  | 212.824875  |   |
| min  | 2.300000e+05 | 27.337938   | 30.140435   | 916.531878  | 916.531878  | 916.531878  |   |
| max  | 1.028000e+06 | 5.425040    | -345.150024 | 1020.616510 | 1020.616510 | 1020.616510 |   |
| 50%  | 7.092500e+04 | 22.500000   | -16.337463  | 1262.712524 | 1262.712524 | 1262.712524 |   |
| 75%  | 1.080000e+05 | 35.750000   | -1.750000   | 1977.000000 | 1977.000000 | 1977.000000 |   |
| max  | 4.395838e+05 | 54.109971   | 13.400024   | 3152.174067 | 3152.174067 | 3152.174067 |   |

|      | count       | 3d MA       | 10d MA      | Price_Rise  | Year        | 3d MA       | % |
|------|-------------|-------------|-------------|-------------|-------------|-------------|---|
| mean | 2415.000000 | 2415.000000 | 2415.000000 | 2415.000000 | 2415.000000 | 2415.000000 |   |
| std  | 2126.355000 | 2126.355000 | 2126.355000 | 2126.355000 | 2126.355000 | 2126.355000 |   |
| min  | 916.531878  | 916.531878  | 916.531878  | 916.531878  | 916.531878  | 916.531878  |   |
| max  | 1057.000000 | 1057.000000 | 1057.000000 | 1057.000000 | 1057.000000 | 1057.000000 |   |

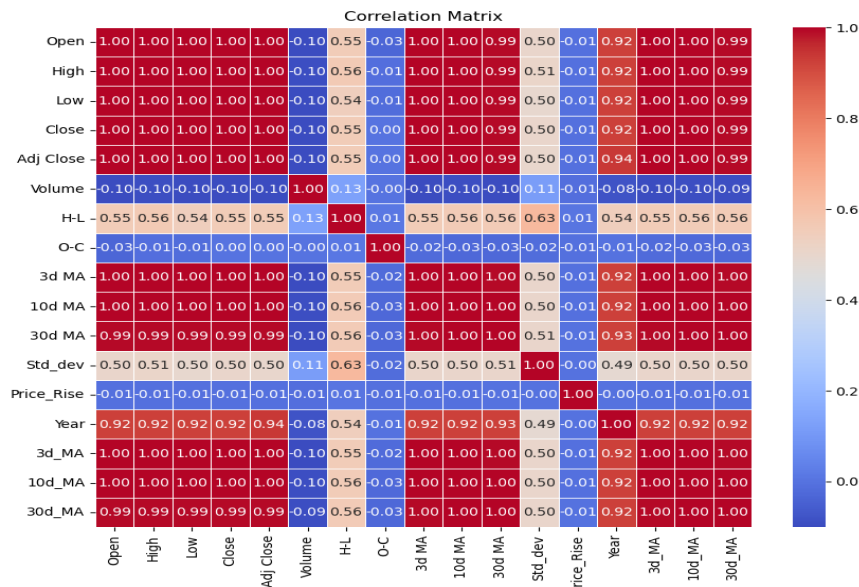
|      | count       | 3d MA       | 10d MA      | Price_Rise  | Year        | 3d MA       | % |
|------|-------------|-------------|-------------|-------------|-------------|-------------|---|
| mean | 2415.000000 | 2415.000000 | 2415.000000 | 2415.000000 | 2415.000000 | 2415.000000 |   |
| std  | 2126.355000 | 2126.355000 | 2126.355000 | 2126.355000 | 2126.355000 | 2126.355000 |   |
| min  | 916.531878  | 916.531878  | 916.531878  | 916.531878  | 916.531878  | 916.531878  |   |
| max  | 1057.000000 | 1057.000000 | 1057.000000 | 1057.000000 | 1057.000000 | 1057.000000 |   |

From the descriptive statistics tables (Fig.6) a total of 2415 rows are observed due to data reduction techniques (previously 2444 rows). The TCS stock dataset, after data refinement, shows a notable growth in mean closing price from 2014 to 2023. Price volatility is evident from the standard deviation, and the nearly balanced 'Price\_Rise' indicates an equal distribution



of gain and loss days. Moving averages smooth out short-term fluctuations, underlining the stock's long-term upward trajectory.

Figure 7: correlation matrix of data variables



The correlation matrix for TCS shows a strong positive correlation among the stock's Open, High, Low, Close, and Adjusted Close prices, indicative of their typical co-movement. A notable negative correlation between 'Year' and price variables may point to a long-term downtrend (Fig.7). 'Volume' and 'Price\_Rise' display weak correlations with price metrics,

## 4.MACHINE LEARNING CLASSIFICATION

Two machine learning models, Logistic Regression and Extra Trees Classifier, were assessed for their efficacy in predicting stock price movements using historical data from TCS. Features ranging from price differentials (H-L) to volatility (Std\_dev) were utilized.

The dataset was split into training and testing sets, with a 20% test size, and feature scaling was applied using StandardScaler for enhanced model performance evaluation. This analysis aimed to determine the most effective model for predicting stock price trends based on historical data.

### A.LOGISTIC REGRESSION

Logistic Regression, a statistical method for binary classification, applies a logistic function to estimate the probability of an event based on inputs. ( David G. Kleinbaum , Mitchel Klein, 2010) When predicting stock movements, such as for TCS, it evaluates linear relationships between features like price range and volatility to forecast price trends, as detailed in section 6.

Formula:

$$P(y = 1|x) = \frac{1}{1+e^{-(\beta_0+\beta_1x)}}$$

Where:

- $P(y = 1|x)$  is the probability of the positive class.
- $\beta_0$  is the intercept.
- $\beta_1$  is the coefficient for the feature  $x$ .

## B. EXTRA TREES

Extra Trees, an ensemble of decision trees, randomizes split points and features to build diverse trees, reducing overfitting and increasing robustness to noise. This method, faster than many tree-based algorithms, excels in both classification and regression tasks (Yazan Ahmad Alsariera; Victor Elijah Adeyemo; Abdullateef Oluwagbemiga Balogun, 2020), as highlighted in section 6.

## 5.CROSS VALIDATION ACCURACY ANALYSIS

Cross-validation rigorously tests a model's ability to generalize by rotating through different training and testing data subsets, thus ensuring the model's reliability, and guiding the optimization of its settings.

Formula:

$$\text{Mean Accuracy} = \frac{1}{k} \sum_{i=1}^k \text{Accuracy}_i$$

where  $\text{Accuracy}_i$  is the accuracy of the model on the  $i$ -th fold.

$$\text{Standard Deviation} = \sqrt{\frac{1}{k} \sum_{i=1}^k (\text{Accuracy}_i - \text{Mean Accuracy})^2}$$

## A. CROSS VALIDATION - LOGISTIC REGRESSION

The logistic regression model, as evaluated through 5-fold cross-validation, exhibits an average accuracy of 49%. This performance level indicates that the model performs marginally better

than random chance when generalized to the entire dataset, and the low standard deviation signifies stable performance across various data splits.

## B. CROSS VALIDATION - EXTRA TREES

The Extra Trees Classifier, upon evaluation with cross-validation, yielded an average accuracy of 48%. This result suggests that the model's predictive capability is nearly equivalent to random guessing. Its performance indicates low variability, suggesting consistent model behaviour across various subsets of the dataset.

## 6.EVALUATION OF CLASSIFIERS RESULTS

The Logistic Regression model, with an overall accuracy of 49%, showed proficiency in predicting stock price increases with a 49% precision and 88% recall, but it was less effective at predicting when prices would not rise, with a precision of 46% and a recall of only 10%. The model's f1-scores reflect this imbalance, suggesting it's better at forecasting price rises than declines.

|              | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0            | 0.46      | 0.10   | 0.17     | 244     |
| 1            | 0.49      | 0.88   | 0.63     | 239     |
| accuracy     |           |        | 0.49     | 483     |
| macro avg    | 0.48      | 0.49   | 0.40     | 483     |
| weighted avg | 0.48      | 0.49   | 0.40     | 483     |

he Extra Trees Classifier achieved a balanced performance with an overall accuracy of 48%, showing near-equal precision in predicting both 'no rise' and 'price rise' outcomes at around 49% and 48% respectively. It was better at identifying 'price rise' cases with a 59% success rate, compared to 38% for 'no rise' scenarios. The f1-scores suggest moderate effectiveness across predictions, highlighting its consistent but modest ability to forecast stock movements.

| Classification Report: |           |        |          |         |  |
|------------------------|-----------|--------|----------|---------|--|
|                        | precision | recall | f1-score | support |  |
| 0                      | 0.49      | 0.38   | 0.42     | 244     |  |
| 1                      | 0.48      | 0.59   | 0.53     | 239     |  |
| accuracy               |           |        | 0.48     | 483     |  |
| macro avg              | 0.48      | 0.49   | 0.48     | 483     |  |
| weighted avg           | 0.48      | 0.48   | 0.48     | 483     |  |

When applied to a 20% dataset split, both classifiers showed similar performance. Logistic regression excelled in predicting price rises, indicated by higher recall and f1-score for this outcome, while the Extra Trees Classifier delivered more balanced results across scenarios, with comparable precision and recall for both 'no rise' and 'rise' predictions.

```
In [46]: from sklearn.model_selection import cross_val_score
from sklearn.metrics import make_scorer, accuracy_score

# Perform cross-validation
accuracy_scores = cross_val_score(model_lr, X, Y, cv=5, scoring=make_scorer(accuracy_score))

# Print mean and standard deviation of accuracy
print(f"Mean Accuracy: {accuracy_scores.mean():.2f}")
print(f"Standard Deviation: {accuracy_scores.std():.2f}")

Mean Accuracy: 0.49
Standard Deviation: 0.02
```

```
In [47]: from sklearn.ensemble import ExtraTreesClassifier
from sklearn.model_selection import cross_val_score
from sklearn.metrics import make_scorer, accuracy_score

# Create an Extra Trees Classifier model
model_et = ExtraTreesClassifier()

# Perform cross-validation
accuracy_scores = cross_val_score(model_et, X, Y, cv=5, scoring=make_scorer(accuracy_score))

# Print mean and standard deviation of accuracy
print(f"Mean Accuracy: {accuracy_scores.mean():.2f}")
print(f"Standard Deviation Accuracy: {accuracy_scores.std():.2f}")

Mean Accuracy: 0.50
Standard Deviation Accuracy: 0.02
```

## 7.PREDICTION OF PRICE USING X\_ TEST DATA: EXTRA TREE

The choice to proceed with the Extra Trees model, as discussed in section 5, is justified by its superior performance on the split dataset, where it surpassed the logistic regression model in effectiveness. (by Matthew F. Dixon; Igor Halperin; Paul A. Biloko, 2020)

### A. CLASSIFICATION REPORT


The classification report evaluates model performance on binary classification through precision (accuracy of class predictions), recall (proportion of actual class instances correctly identified), F1-score (balance between precision and recall), and support (instances per class), offering a comprehensive view of effectiveness in classifying instances. (Buchanan, Bonnie G, 2019)

|                        |           |        |          |         |
|------------------------|-----------|--------|----------|---------|
| Classification Report: |           |        |          |         |
|                        | precision | recall | f1-score | support |
| 0                      | 0.49      | 0.38   | 0.42     | 244     |
| 1                      | 0.48      | 0.59   | 0.53     | 239     |
| accuracy               |           |        | 0.48     | 483     |
| macro avg              | 0.48      | 0.49   | 0.48     | 483     |
| weighted avg           | 0.48      | 0.48   | 0.48     | 483     |

## B. CONFUSION MATRIX

The confusion matrix shows the model's prediction accuracy: 92 correct for 'no price rise' and 142 for 'price rise', but with 152 and 97 mispredictions for each category, respectively. It highlights the model's stronger performance in identifying 'price rise' scenarios. (Fig.8).

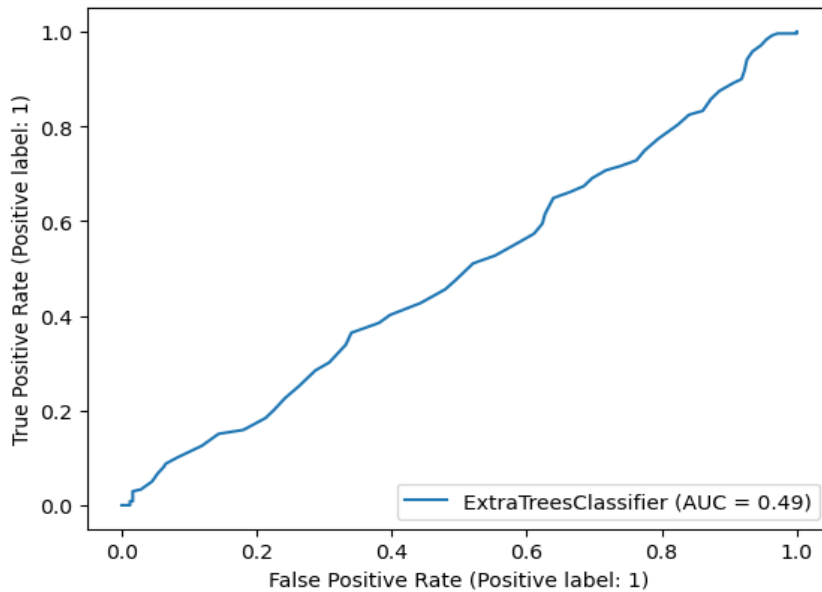
Figure 8: Confusion Matrix of model (Extra Tree model)



## C. ROC CURVE

The ROC curve for the Extra Trees Classifier, with an AUC of 0.49, demonstrates the model's balance between sensitivity and specificity, suggesting its performance in predicting 'price rise' is like random guessing. (Fig.9).

Figure 9: ROC Curve of model (Extra Tree classification model)



## 8. MARKET AND STRATEGY RETURNS

"Market Returns" represent hypothetical earnings from holding a stock, while "Strategy Returns" are based on trading strategies informed by machine learning predictions, such as those from the Extra Trees model. In our analysis, model forecasts are added to a 'Y\_predicted' column to calculate these returns. "Tomorrows Returns" tracks daily logarithmic returns, adjusting for the natural fluctuation between consecutive days' closing prices, ensuring alignment for strategic analysis. (Yumei Yao, 2024)


$$\text{Market Return} = \frac{\text{Closing Price}_{\text{today}} - \text{Closing Price}_{\text{yesterday}}}{\text{Closing Price}_{\text{yesterday}}}$$

Strategy returns are derived by adopting long positions for true 'Y\_predicted' and short positions for false. The 'Strategy Returns' column, initialized at zero, records 'Tomorrows Returns'—positively for predicted price rises and negatively for predicted drops, simulating buying or selling the stock accordingly. (Yumei Yao, 2024)

## 9.CUMMULATIVE RETURNS ANALYSIS

Cumulative return analysis evaluates investment strategy effectiveness, asset performance, and decision profitability, guiding long-term investment insights and strategy adjustments. It underpins trading strategy assessments, with market and strategy cumulative returns calculated using the cumsum() function.

**Figure 10: cumulative market and strategy return based on Y\_ Prediction**



The graph tracks cumulative returns from early 2022 to early 2024, contrasting market performance with a predictive strategy's returns. Initially, market returns outshine the strategy, but from April 2022 onwards, the two approaches' performances begin to align and occasionally switch leads. The strategy notably outperforms the market amid volatility from late 2022 to 2023, suggesting its potential superiority during unstable market conditions. By the period's end, the strategy notably rebounds above market returns, indicating its potential advantage over standard market-following approaches in turbulent times.

## 10.INTERPRETATION AND DISCUSSION

Analysing TCS stock prices using the Extra Trees Classifier revealed a model accuracy of around 48%, which suggests that while the model has some predictive ability, it's not highly reliable for forecasting stock price movements. This outcome aligns with the efficient market hypothesis, which holds that stocks reflect all available information, making prediction challenging. The model's performance underscores the complex and often unpredictable nature

of financial markets, where numerous variables can influence stock prices beyond historical data.

## 11.CONCLUSION

The analysis of TCS stock with the Extra Trees Classifier indicates that machine learning can capture certain stock market trends to an extent, but the model's predictive power is limited. The performance is reflective of market complexities and unpredictability, highlighting the difficulty in forecasting prices with high accuracy due to unexpected market events and the challenge of overfitting. Although machine learning adds analytical value, it should be complemented by traditional financial analysis due to the volatile nature of the markets and the limitations of existing models. For future advancements, research should aim at enhancing the adaptability and robustness of predictive models under varied market conditions.

## 12.REFERENCES

- David G. Kleinbaum , Mitchel Klein. (2010). *Logistic regression*. USA: Springer New York, NY.
- Buchanan, Bonnie G. (2019). *Artificial intelligence in finance*. Finland: The Alan Turing Institute.
- by Matthew F. Dixon; Igor Halperin; Paul A. Biloko. (2020). *Machine learning in finance: from theory to practice*. UK: Springer International Publishing.
- IBM. (2024, 03 24). *EDA*. Retrieved from IBM: <https://www.ibm.com/topics/exploratory-data-analysis>
- Tata Company. (2024, march 24). *TATA Companies*. Retrieved from Tata Consultancy Services: <https://www.tata.com/business/tcs>
- Yazan Ahmad Alsariera; Victor Elijah Adeyemo; Abdullateef Oluwagbemiga Balogun. (2020). *AI Meta-Learners and Extra-Trees Algorithm*. NY: IEEE.
- Yumei Yao. (2024, March 24). Seminar 6. *AI & ML*. London, UK, UK: university of westminster.



## 13.APPENDICES

### 1.Imported Library

```
In [4]: import warnings
warnings.simplefilter(action='ignore', category=FutureWarning)

In [5]: #Import libraries and metrics
import yfinance as yf
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split
from sklearn.model_selection import cross_val_score
from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import ExtraTreesClassifier
from sklearn.metrics import accuracy_score, roc_curve, classification_report
from sklearn.metrics import confusion_matrix, make_scorer, roc_auc_score
from sklearn.metrics import RocCurveDisplay, ConfusionMatrixDisplay
```

### 2.Historical Data (TCS Uploaded) & Cleaned

```
In [ ]: pd.options.mode.chained_assignment = None

In [13]: import pandas as pd
file_path = 'TCS.BO.csv'
df = pd.read_csv(file_path)
print(df.head())
```

|   | Date       | Open        | High        | Low         | Close       | Adj Close \ |
|---|------------|-------------|-------------|-------------|-------------|-------------|
| 0 | 2014-01-01 | 1090.000000 | 1092.500000 | 1075.525024 | 1076.849976 | 858.149536  |
| 1 | 2014-01-02 | 1080.000000 | 1093.574951 | 1075.000000 | 1081.150024 | 861.576538  |
| 2 | 2014-01-03 | 1087.500000 | 1114.125000 | 1075.500000 | 1111.000000 | 885.364075  |
| 3 | 2014-01-06 | 1113.000000 | 1121.824951 | 1098.500000 | 1119.800049 | 892.376953  |
| 4 | 2014-01-07 | 1119.150024 | 1127.400024 | 1100.300049 | 1104.099976 | 879.865173  |

```
Volume
0    73300.0
1   117324.0
2   267204.0
3   145172.0
4   135104.0
```

```
In [19]: df_cleaned = df.dropna()
print(df_cleaned)
```

|      | Date       | Open        | High        | Low         | Close \     |
|------|------------|-------------|-------------|-------------|-------------|
| 0    | 2014-01-01 | 1090.000000 | 1092.500000 | 1075.525024 | 1076.849976 |
| 1    | 2014-01-02 | 1080.000000 | 1093.574951 | 1075.000000 | 1081.150024 |
| 2    | 2014-01-03 | 1087.500000 | 1114.125000 | 1075.500000 | 1111.000000 |
| 3    | 2014-01-06 | 1113.000000 | 1121.824951 | 1098.500000 | 1119.800049 |
| 4    | 2014-01-07 | 1119.150024 | 1127.400024 | 1100.300049 | 1104.099976 |
| ...  | ...        | ...         | ...         | ...         | ...         |
| 2462 | 2023-12-22 | 3796.300049 | 3845.949951 | 3762.699951 | 3825.300049 |
| 2463 | 2023-12-26 | 3780.100098 | 3833.850098 | 3780.100098 | 3794.600098 |
| 2464 | 2023-12-27 | 3795.550049 | 3818.000000 | 3768.100098 | 3810.800049 |
| 2465 | 2023-12-28 | 3822.100098 | 3838.250000 | 3793.750000 | 3801.050049 |
| 2466 | 2023-12-29 | 3797.850098 | 3822.949951 | 3766.050049 | 3794.949951 |

|      | Adj Close   | Volume   |
|------|-------------|----------|
| 0    | 858.149536  | 73300.0  |
| 1    | 861.576538  | 117324.0 |
| 2    | 885.364075  | 267204.0 |
| 3    | 892.376953  | 145172.0 |
| 4    | 879.865173  | 135104.0 |
| ...  | ...         | ...      |
| 2462 | 3825.300049 | 127163.0 |
| 2463 | 3794.600098 | 70216.0  |
| 2464 | 3810.800049 | 28290.0  |
| 2465 | 3801.050049 | 29256.0  |
| 2466 | 3794.949951 | 105711.0 |

### 3.Feature Engineering

```
In [20]: # Calculate 'H-L' and 'O-C' differences
df_cleaned['H-L'] = df_cleaned['High'] - df_cleaned['Low']
df_cleaned['O-C'] = df_cleaned['Close'] - df_cleaned['Open']

# Calculate moving averages
df_cleaned['3d MA'] = df_cleaned['Close'].rolling(window=3).mean()
df_cleaned['10d MA'] = df_cleaned['Close'].rolling(window=10).mean()
df_cleaned['30d MA'] = df_cleaned['Close'].rolling(window=30).mean()

# Calculate standard deviation
df_cleaned['Std_dev'] = df_cleaned['Close'].rolling(window=5).std()

# Create 'Price_Rise' column
df_cleaned['Price_Rise'] = (df_cleaned['Close'].shift(-1) > df_cleaned['Close']).astype(int)

# Drop Last row as it will have NaN values for 'Price_Rise'
df_cleaned = df_cleaned.dropna()

# Display the cleaned DataFrame
print(df_cleaned)
```

|      | Date        | Open        | High        | Low         | Close       | \           |   |
|------|-------------|-------------|-------------|-------------|-------------|-------------|---|
| 29   | 2014-02-11  | 1047.500000 | 1063.000000 | 1047.500000 | 1050.500000 |             |   |
| 30   | 2014-02-12  | 1061.849976 | 1061.900024 | 1047.550049 | 1051.900024 |             |   |
| 31   | 2014-02-13  | 1054.500000 | 1069.500000 | 1049.500000 | 1066.875000 |             |   |
| 32   | 2014-02-14  | 1069.500000 | 1092.199951 | 1064.250000 | 1083.125000 |             |   |
| 33   | 2014-02-17  | 1084.500000 | 1093.500000 | 1074.775024 | 1084.300049 |             |   |
| ...  | ...         | ...         | ...         | ...         | ...         |             |   |
| 2462 | 2023-12-22  | 3796.300049 | 3845.949951 | 3762.699951 | 3825.300049 |             |   |
| 2463 | 2023-12-26  | 3780.100098 | 3833.850098 | 3780.100098 | 3794.600098 |             |   |
| 2464 | 2023-12-27  | 3795.550049 | 3818.000000 | 3768.100098 | 3810.800049 |             |   |
| 2465 | 2023-12-28  | 3822.100098 | 3838.250000 | 3793.750000 | 3801.050049 |             |   |
| 2466 | 2023-12-29  | 3797.850098 | 3822.949951 | 3766.050049 | 3794.949951 |             |   |
|      | Adj Close   | Volume      | H-L         | O-C         | 3d MA       | 10d MA      | \ |
| 29   | 840.140564  | 202536.0    | 15.500000   | 3.000000    | 1056.341675 | 1085.967493 |   |
| 30   | 841.260193  | 126528.0    | 14.349975   | -9.949952   | 1049.750000 | 1080.679993 |   |
| 31   | 853.236633  | 165496.0    | 20.000000   | 12.375000   | 1056.425008 | 1076.472498 |   |
| 32   | 866.232483  | 149938.0    | 27.949951   | 13.625000   | 1067.300008 | 1072.975000 |   |
| 33   | 867.172363  | 202532.0    | 18.724976   | -0.199951   | 1078.100016 | 1071.572510 |   |
| ...  | ...         | ...         | ...         | ...         | ...         | ...         |   |
| 2462 | 3825.300049 | 127163.0    | 83.250000   | 29.000000   | 3798.366699 | 3749.985034 |   |
| 2463 | 3794.600098 | 70216.0     | 53.750000   | 14.500000   | 3803.116699 | 3765.225049 |   |
| 2464 | 3810.800049 | 28290.0     | 49.899902   | 15.250000   | 3810.233399 | 3779.100049 |   |
| 2465 | 3801.050049 | 29256.0     | 44.500000   | -21.050049  | 3802.150065 | 3799.830054 |   |
| 2466 | 3794.949951 | 105711.0    | 56.899902   | -2.900147   | 3802.266683 | 3812.665039 |   |
|      | 30d MA      | Std_dev     | Price_Rise  |             |             |             |   |
| 29   | 1115.374162 | 21.631318   | 1           |             |             |             |   |
| 30   | 1114.542497 | 17.325573   | 1           |             |             |             |   |
| 31   | 1114.066663 | 10.983936   | 1           |             |             |             |   |
| 32   | 1113.137496 | 15.088916   | 1           |             |             |             |   |
| 33   | 1111.954163 | 16.271172   | 0           |             |             |             |   |
| ...  | ...         | ...         | ...         |             |             |             |   |
| 2462 | 3577.036678 | 30.812797   | 0           |             |             |             |   |

### 4.EDA of created Features

```
In [26]: import pandas as pd
import matplotlib.pyplot as plt

# Assuming your DataFrame is named df_cleaned

# Create a larger figure and axis objects
fig, ax1 = plt.subplots(figsize=(14, 7))


# Plot Closing Price on the first y-axis
color = 'tab:blue'
ax1.set_ylabel('Closing Price', color=color)
ax1.plot(df_cleaned['Date'], df_cleaned['Close'], label='Close', color=color)
ax1.tick_params(axis='y', labelcolor=color)

# Create a second y-axis sharing the same x-axis
ax2 = ax1.twinx()

# Plot Volume on the second y-axis
color = 'tab:orange'
ax2.set_ylabel('Volume', color=color)
ax2.plot(df_cleaned['Date'], df_cleaned['Volume'], label='Volume', color=color)
ax2.tick_params(axis='y', labelcolor=color)

# Customize x-axis ticks to show only years
years = pd.to_datetime(df_cleaned['Date']).dt.year.unique()
plt.xticks(ticks=df_cleaned['Date'][::len(df_cleaned)//len(years)], labels=years, rotation=45)

# Add title and Legend
plt.title('Closing Price and Volume Relationship')
fig.tight_layout()
plt.show()
```



```
In [33]: import matplotlib.pyplot as plt

# Assuming your DataFrame is named df_cleaned
# Make sure you have already calculated the moving averages

# Extract years from the 'Date' column
years = pd.to_datetime(df_cleaned['Date']).dt.year.unique()

# Get the indices corresponding to the beginning of each year
year_indices = []
for year in years:
    year_indices.append(df_cleaned[df_cleaned['Date'].str.contains(str(year))].index[0])

plt.figure(figsize=(10, 8))

plt.plot(df_cleaned['3d_MA'], label='3-day MA')
plt.plot(df_cleaned['10d_MA'], label='10-day MA')
plt.plot(df_cleaned['30d_MA'], label='30-day MA')

# Set x-axis ticks to correspond to the indices where the years change
plt.xticks(ticks=year_indices, labels=years)

plt.title('Moving Averages Over Time')
plt.ylabel('Moving Averages')
plt.legend()
plt.show()
```



```
In [35]: import matplotlib.pyplot as plt

# Extract years from the 'Date' column
years = pd.to_datetime(df_cleaned['Date']).dt.year.unique()

fig, axes = plt.subplots(1, 3, figsize=(15, 5))


for ax in axes:
    ax.set_xticks(year_indices)
    ax.set_xticklabels(years)

axes[0].plot(df_cleaned['O-C'])
axes[0].set_title('O-C')

axes[1].plot(df_cleaned['H-L'])
axes[1].set_title('H-L')

axes[2].plot(df_cleaned['Std_dev'])
axes[2].set_title('Std_dev')

plt.show()
```



```
In [37]: import seaborn as sns
import matplotlib.pyplot as plt


fig, axes = plt.subplots(1, 3, figsize=(15,3)) # 1 row, 3 columns

sns.histplot(data=df_cleaned, x="3d_MA", kde=True, stat="density", ax=axes[0])
axes[0].set_title('3-day MA')

sns.histplot(data=df_cleaned, x="10d_MA", kde=True, stat="density", ax=axes[1])
axes[1].set_title('10-day MA')

sns.histplot(data=df_cleaned, x="30d_MA", kde=True, stat="density", ax=axes[2])
axes[2].set_title('30-day MA')


plt.show()
```



```
In [38]: import seaborn as sns
import matplotlib.pyplot as plt

# Price Rise & Fall Count Plot
chart_count = sns.FacetGrid(df_cleaned, col='Price_Rise')
chart_count.map(sns.histplot, 'Close')
plt.suptitle('Price Rise & Fall Count Plot in relation to Close', y=1.02)
plt.show()

# Price Rise & Fall Scatter Plot
chart_scatter = sns.FacetGrid(df_cleaned, col='Price_Rise')
chart_scatter.map(plt.scatter, 'Close', 'Std_dev')
plt.suptitle('Price Rise & Fall Scatter Plot in relation to Close and Std_dev', y=1.02)
plt.show()
```

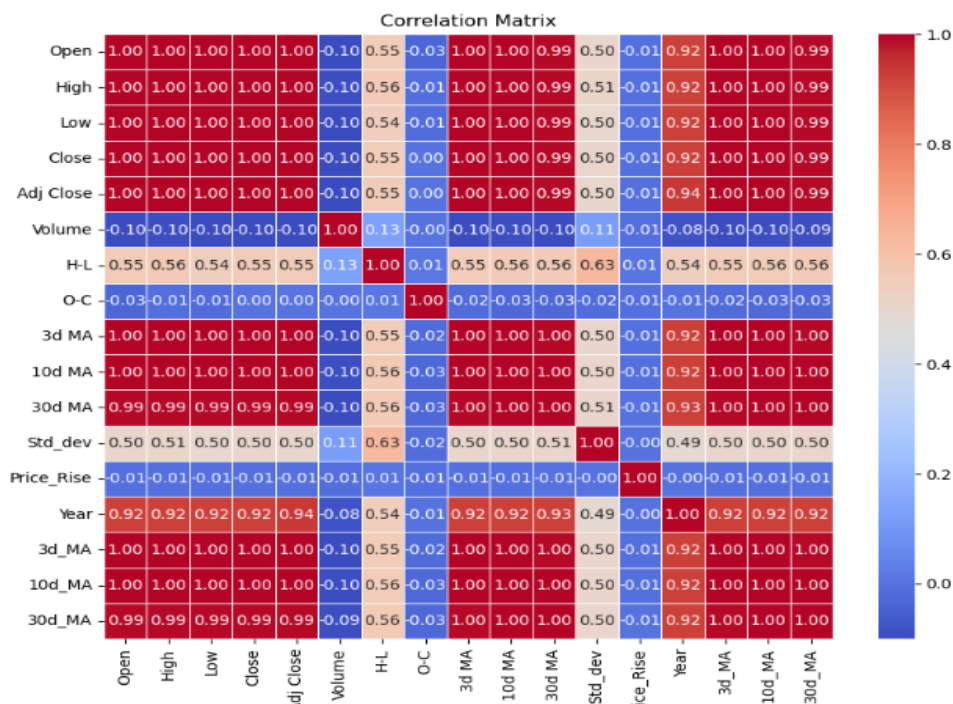


## 5. Correlation Matrix

```
In [41]: import seaborn as sns
import matplotlib.pyplot as plt

# Calculate correlation matrix excluding non-numeric columns
corr_matrix = df_cleaned.corr(method='pearson', min_periods=1, numeric_only=True)

# Plot correlation matrix
plt.figure(figsize=(10, 8))
sns.heatmap(corr_matrix, annot=True, cmap='coolwarm', fmt=".2f", linewidths=.5)
plt.title('Correlation Matrix')
plt.show()
```



## 6. Train and Test Data Splitting

```
In [43]: from sklearn.model_selection import train_test_split

# Split the data into training and testing sets
X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size=0.2, shuffle=False)

# Display the testing set features
print(X_test)
```

|      | H-L       | O-C        | 3d MA       | 10d MA      | 30d MA      | Std_dev   |
|------|-----------|------------|-------------|-------------|-------------|-----------|
| 1967 | 40.949951 | 22.300049  | 3633.750000 | 3606.595044 | 3556.254997 | 40.325331 |
| 1968 | 58.599853 | -5.750000  | 3654.350016 | 3609.965039 | 3562.326660 | 45.779676 |
| 1969 | 46.850098 | 25.550049  | 3676.200033 | 3618.510034 | 3568.028328 | 34.340124 |
| 1970 | 30.449951 | -7.699951  | 3691.033366 | 3626.960034 | 3573.141667 | 30.188729 |
| 1971 | 35.050049 | -6.850098  | 3698.666667 | 3639.315015 | 3577.741667 | 18.588547 |
| ...  | ...       | ...        | ...         | ...         | ...         | ...       |
| 2462 | 83.250000 | 29.000000  | 3798.366699 | 3749.985034 | 3577.036678 | 30.812797 |
| 2463 | 53.750000 | 14.500000  | 3803.116699 | 3765.225049 | 3591.915015 | 18.113791 |
| 2464 | 49.899902 | 15.250000  | 3810.233399 | 3779.100049 | 3607.848348 | 17.911341 |
| 2465 | 44.500000 | -21.050049 | 3802.150065 | 3799.830054 | 3623.493351 | 14.216267 |
| 2466 | 56.899902 | -2.900147  | 3802.266683 | 3812.665039 | 3636.681681 | 12.941387 |

[483 rows x 6 columns]

## 7. Logistic regression

```
In [45]: from sklearn.linear_model import LogisticRegression
from sklearn.metrics import classification_report

# Create a Logistic Regression model
model_lr = LogisticRegression(random_state=101)

# Train the model
model_lr.fit(X_train, Y_train)

# Predict on the test set
Y_pred = model_lr.predict(X_test)

# Print the classification report
print(classification_report(Y_test, Y_pred))
```

|              | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0            | 0.46      | 0.10   | 0.17     | 244     |
| 1            | 0.49      | 0.88   | 0.63     | 239     |
| accuracy     |           |        | 0.49     | 483     |
| macro avg    | 0.48      | 0.49   | 0.40     | 483     |
| weighted avg | 0.48      | 0.49   | 0.40     | 483     |

## 8. Extra Tree classification

```
In [48]: from sklearn.ensemble import ExtraTreesClassifier
from sklearn.metrics import classification_report

# Create an Extra Trees Classifier model
model_et = ExtraTreesClassifier(random_state=101)

# Train the model
model_et.fit(X_train, Y_train)

# Predict on the test set
Y_pred = model_et.predict(X_test)

# Print classification report
classification_rep = classification_report(Y_test, Y_pred)
print("Classification Report:\n", classification_rep)
```

Classification Report:

|              | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0            | 0.49      | 0.38   | 0.42     | 244     |
| 1            | 0.48      | 0.59   | 0.53     | 239     |
| accuracy     |           |        | 0.48     | 483     |
| macro avg    | 0.48      | 0.49   | 0.48     | 483     |
| weighted avg | 0.48      | 0.48   | 0.48     | 483     |

## 9. Cross Validation

```
In [46]: from sklearn.model_selection import cross_val_score
from sklearn.metrics import make_scorer, accuracy_score

# Perform cross-validation
accuracy_scores = cross_val_score(model_lr, X, Y, cv=5, scoring=make_scorer(accuracy_score))

# Print mean and standard deviation of accuracy
print(f"Mean Accuracy: {accuracy_scores.mean():.2f}")
print(f"Standard Deviation: {accuracy_scores.std():.2f}")
```

Mean Accuracy: 0.49  
Standard Deviation: 0.02

```
In [47]: from sklearn.ensemble import ExtraTreesClassifier
from sklearn.model_selection import cross_val_score
from sklearn.metrics import make_scorer, accuracy_score

# Create an Extra Trees Classifier model
model_et = ExtraTreesClassifier()

# Perform cross-validation
accuracy_scores = cross_val_score(model_et, X, Y, cv=5, scoring=make_scorer(accuracy_score))

# Print mean and standard deviation of accuracy
print(f"Mean Accuracy: {accuracy_scores.mean():.2f}")
print(f"Standard Deviation Accuracy: {accuracy_scores.std():.2f}")
```

Mean Accuracy: 0.50  
Standard Deviation Accuracy: 0.02

## 10. Prediction of Price Rise Using extra Tree on X test Data

```
In [48]: from sklearn.ensemble import ExtraTreesClassifier
from sklearn.metrics import classification_report

# Create an Extra Trees Classifier model
model_et = ExtraTreesClassifier(random_state=101)

# Train the model
model_et.fit(X_train, Y_train)

# Predict on the test set
Y_pred = model_et.predict(X_test)

# Print classification report
classification_rep = classification_report(Y_test, Y_pred)
print("Classification Report:\n", classification_rep)
```

Classification Report:


|              | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0            | 0.49      | 0.38   | 0.42     | 244     |
| 1            | 0.48      | 0.59   | 0.53     | 239     |
| accuracy     |           |        | 0.48     | 483     |
| macro avg    | 0.48      | 0.49   | 0.48     | 483     |
| weighted avg | 0.48      | 0.48   | 0.48     | 483     |

## 11. Confusion Matrix

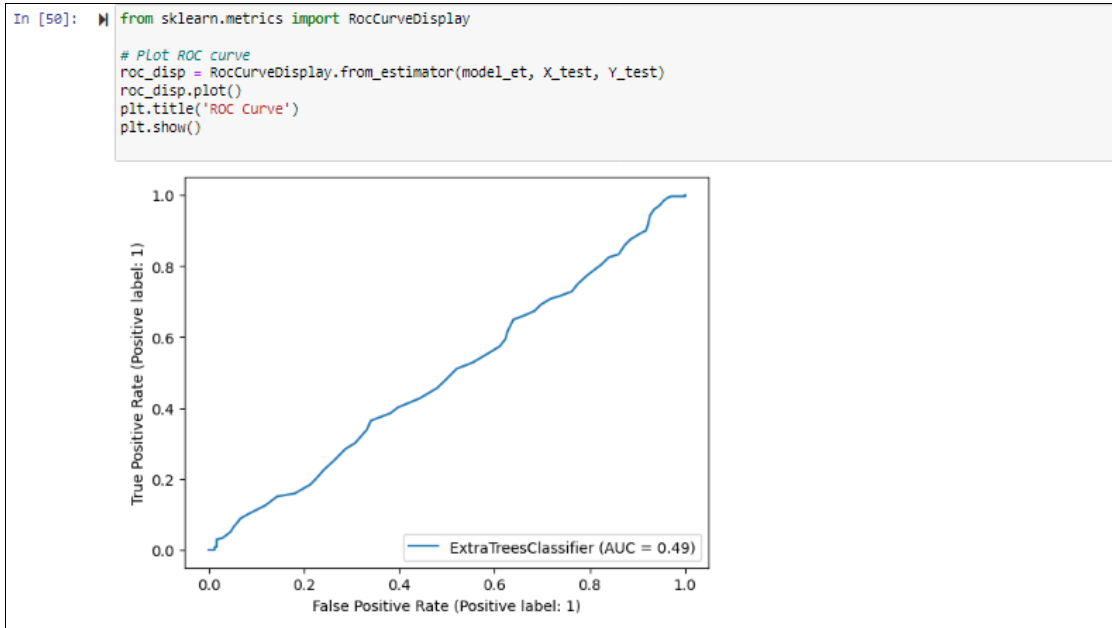
```
In [49]: from sklearn.metrics import confusion_matrix, ConfusionMatrixDisplay

# Calculate the confusion matrix
conf_matrix = confusion_matrix(Y_test, Y_pred)

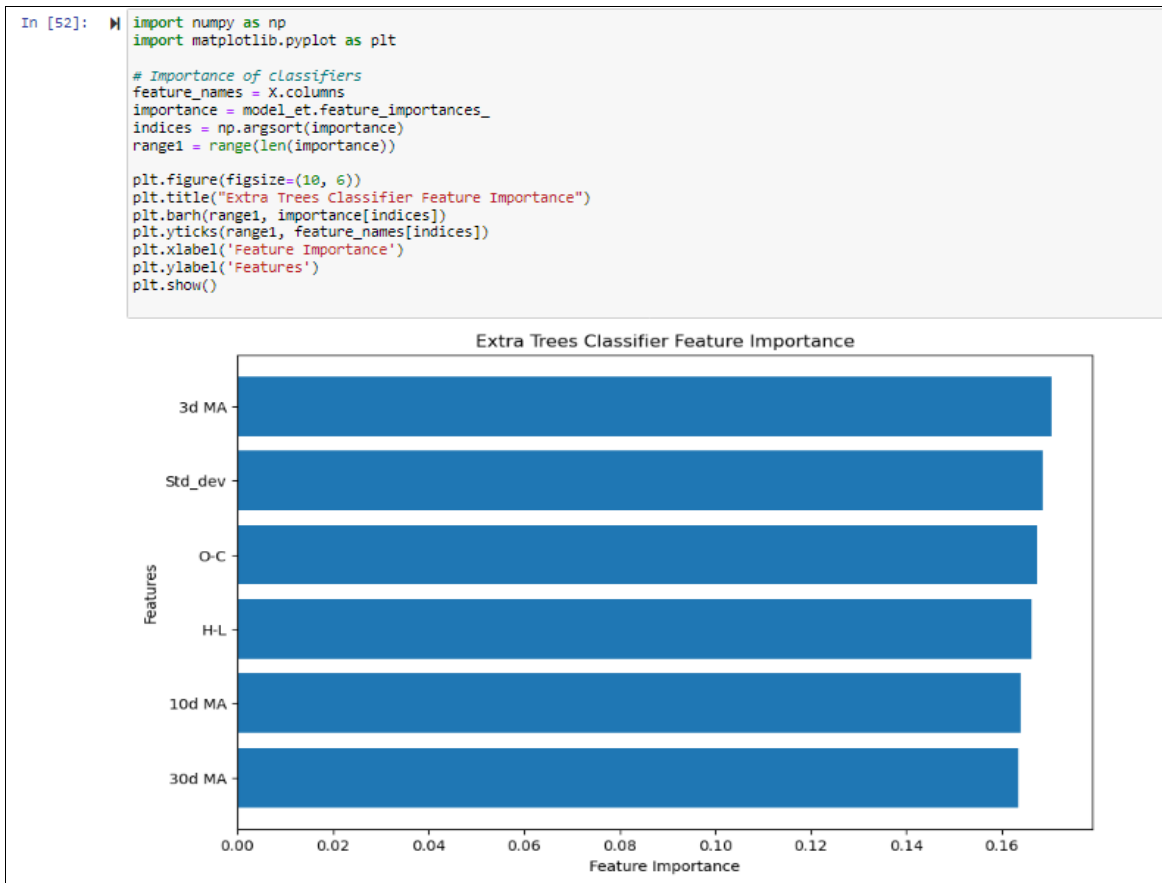
# Display the confusion matrix
disp = ConfusionMatrixDisplay(confusion_matrix=conf_matrix, display_labels=model_et.classes_)
disp.plot(cmap='Blues')
plt.title('Confusion Matrix')
plt.show()
```



## 12.ROC Curve of the model (Extra Tree)



## 13.Importance of Classifier in model





## 14. Market and Strategic Return

```
In [55]: # Create a new column for market returns
df_cleaned['Market Returns'] = 0.0

# Calculate market returns
df_cleaned['Market Returns'] = np.log(df_cleaned['Close'] / df_cleaned['Close'].shift(1))

# Shift market returns by one day to represent tomorrow's returns
df_cleaned['Market Returns'] = df_cleaned['Market Returns'].shift(-1)

# Drop any rows with NaN values
df_cleaned.dropna(inplace=True)

# Display the DataFrame
df_cleaned
```

Out[55]:

|      | Date       | Open        | High        | Low         | Close       | Adj Close   | Volume   | H-L       | O-C        | 3d MA       | 10d MA      | 30d MA      |
|------|------------|-------------|-------------|-------------|-------------|-------------|----------|-----------|------------|-------------|-------------|-------------|
| 1967 | 2021-12-23 | 3840.800000 | 3871.000000 | 3830.050049 | 3862.800049 | 3538.583498 | 82978.0  | 40.949951 | 22.300049  | 3833.750000 | 3808.595044 | 3588.254997 |
| 1968 | 2021-12-24 | 3878.000000 | 3703.899902 | 3845.300049 | 3870.250000 | 3543.776855 | 141802.0 | 58.599853 | -5.750000  | 3854.350016 | 3809.985039 | 3682.328880 |
| 1969 | 2021-12-27 | 3870.000000 | 3700.000000 | 3853.149902 | 3895.550049 | 3588.204834 | 37807.0  | 48.850098 | 25.550049  | 3878.200033 | 3818.510034 | 3688.028328 |
| 1970 | 2021-12-28 | 3715.000000 | 3724.500000 | 3894.050049 | 3707.300049 | 3579.550049 | 34438.0  | 30.449951 | -7.899951  | 3891.033388 | 3828.980034 | 3573.141887 |
| 1971 | 2021-12-29 | 3700.000000 | 3719.449951 | 3884.399902 | 3893.149902 | 3585.887207 | 75985.0  | 35.050049 | -8.850098  | 3898.888887 | 3839.315015 | 3577.741887 |
| ...  | ...        | ...         | ...         | ...         | ...         | ...         | ...      | ...       | ...        | ...         | ...         | ...         |
| 2461 | 2023-12-21 | 3734.399902 | 3805.000000 | 3734.399902 | 3789.449951 | 3789.449951 | 90558.0  | 70.600098 | 55.050049  | 3793.988718 | 3730.095020 | 3582.218343 |
| 2462 | 2023-12-22 | 3798.300049 | 3845.949951 | 3782.899951 | 3825.300049 | 3825.300049 | 127183.0 | 83.250000 | 29.000000  | 3798.388889 | 3749.985034 | 3577.038878 |
| 2463 | 2023-12-26 | 3780.100098 | 3833.850098 | 3780.100098 | 3794.800098 | 3794.800098 | 70218.0  | 53.750000 | 14.500000  | 3803.118889 | 3785.225049 | 3591.915015 |
| 2464 | 2023-12-27 | 3795.550049 | 3818.000000 | 3788.100098 | 3810.800049 | 3810.800049 | 28290.0  | 49.899902 | 15.250000  | 3810.233399 | 3779.100049 | 3807.848348 |
| 2465 | 2023-12-28 | 3822.100098 | 3838.250000 | 3793.750000 | 3801.050049 | 3801.050049 | 29258.0  | 44.500000 | -21.050049 | 3802.150085 | 3799.830054 | 3823.493351 |

482 rows x 20 columns

```
In [56]: # Create a new column for strategy returns
df_cleaned['Strategy Returns'] = 0.0

# Compute strategy returns based on Y_pred
df_cleaned['Strategy Returns'] = np.where(df_cleaned['Y_pred'] == True,
                                          df_cleaned['Market Returns'],
                                          -df_cleaned['Market Returns'])

# Display the DataFrame
df_cleaned
```

Out[56]:

|      | Date       | Open        | High        | Low         | Close       | Adj Close   | Volume   | H-L       | O-C        | 3d MA       | ... | 30d MA      | Std_dev   | f   |
|------|------------|-------------|-------------|-------------|-------------|-------------|----------|-----------|------------|-------------|-----|-------------|-----------|-----|
| 1967 | 2021-12-23 | 3840.800000 | 3871.000000 | 3830.050049 | 3862.800049 | 3538.583498 | 82978.0  | 40.949951 | 22.300049  | 3833.750000 | ... | 3588.254997 | 40.325331 |     |
| 1968 | 2021-12-24 | 3878.000000 | 3703.899902 | 3845.300049 | 3870.250000 | 3543.776855 | 141802.0 | 58.599853 | -5.750000  | 3854.350016 | ... | 3682.328880 | 45.779878 |     |
| 1969 | 2021-12-27 | 3870.000000 | 3700.000000 | 3853.149902 | 3895.550049 | 3588.204834 | 37807.0  | 48.850098 | 25.550049  | 3878.200033 | ... | 3688.028328 | 34.340124 |     |
| 1970 | 2021-12-28 | 3715.000000 | 3724.500000 | 3894.050049 | 3707.300049 | 3579.550049 | 34438.0  | 30.449951 | -7.899951  | 3891.033388 | ... | 3573.141887 | 30.188729 |     |
| 1971 | 2021-12-29 | 3700.000000 | 3719.449951 | 3884.399902 | 3893.149902 | 3585.887207 | 75985.0  | 35.050049 | -8.850098  | 3898.888887 | ... | 3577.741887 | 18.588547 |     |
| ...  | ...        | ...         | ...         | ...         | ...         | ...         | ...      | ...       | ...        | ...         | ... | ...         | ...       | ... |
| 2461 | 2023-12-21 | 3734.399902 | 3805.000000 | 3734.399902 | 3789.449951 | 3789.449951 | 90558.0  | 70.600098 | 55.050049  | 3793.988718 | ... | 3582.218343 | 37.489438 |     |
| 2462 | 2023-12-22 | 3798.300049 | 3845.949951 | 3782.899951 | 3825.300049 | 3825.300049 | 127183.0 | 83.250000 | 29.000000  | 3798.388889 | ... | 3577.038878 | 30.812797 |     |
| 2463 | 2023-12-26 | 3780.100098 | 3833.850098 | 3780.100098 | 3794.800098 | 3794.800098 | 70218.0  | 53.750000 | 14.500000  | 3803.118889 | ... | 3591.915015 | 18.113791 |     |
| 2464 | 2023-12-27 | 3795.550049 | 3818.000000 | 3788.100098 | 3810.800049 | 3810.800049 | 28290.0  | 49.899902 | 15.250000  | 3810.233399 | ... | 3807.848348 | 17.911341 |     |
| 2465 | 2023-12-28 | 3822.100098 | 3838.250000 | 3793.750000 | 3801.050049 | 3801.050049 | 29258.0  | 44.500000 | -21.050049 | 3802.150085 | ... | 3823.493351 | 14.218287 |     |

482 rows x 21 columns

## 15. cumulative market and strategic return

```
In [57]: # Compute cumulative market returns
df_cleaned['Cumulative Market Returns'] = np.cumsum(df_cleaned['Market Returns'])

# Compute cumulative strategy returns
df_cleaned['Cumulative Strategy Returns'] = np.cumsum(df_cleaned['Strategy Returns'])

# Display the DataFrame
df_cleaned
```

Out[57]:

|      | Date       | Open        | High        | Low         | Close       | Adj Close   | Volume   | H-L       | O-C        | 3d MA       | ... | Price_Rise | Year | 3c      |
|------|------------|-------------|-------------|-------------|-------------|-------------|----------|-----------|------------|-------------|-----|------------|------|---------|
| 1967 | 2021-12-23 | 3640.500000 | 3671.000000 | 3630.050049 | 3662.800049 | 3536.583496 | 82978.0  | 40.949951 | 22.300049  | 3633.750000 | ... | 1          | 2021 | 3633.75 |
| 1968 | 2021-12-24 | 3676.000000 | 3703.899902 | 3645.300049 | 3670.250000 | 3543.776855 | 141802.0 | 58.599853 | -6.750000  | 3654.350016 | ... | 1          | 2021 | 3654.35 |
| 1969 | 2021-12-27 | 3670.000000 | 3700.000000 | 3653.149902 | 3695.550049 | 3568.204834 | 37807.0  | 46.850098 | 25.550049  | 3676.200033 | ... | 1          | 2021 | 3676.20 |
| 1970 | 2021-12-28 | 3715.000000 | 3724.500000 | 3694.050049 | 3707.300049 | 3579.550049 | 34438.0  | 30.449951 | -7.699951  | 3691.033366 | ... | 0          | 2021 | 3691.03 |
| 1971 | 2021-12-29 | 3700.000000 | 3719.449951 | 3684.399902 | 3693.149902 | 3565.887207 | 75985.0  | 35.050049 | -8.850098  | 3698.686867 | ... | 1          | 2021 | 3698.68 |
| ...  | ...        | ...         | ...         | ...         | ...         | ...         | ...      | ...       | ...        | ...         | ... | ...        | ...  | ...     |
| 2461 | 2023-12-21 | 3734.399902 | 3805.000000 | 3734.399902 | 3789.449951 | 3789.449951 | 90556.0  | 70.800098 | 55.050049  | 3793.986716 | ... | 1          | 2023 | 3793.98 |
| 2462 | 2023-12-22 | 3798.300049 | 3845.949951 | 3782.699951 | 3825.300049 | 3825.300049 | 127163.0 | 83.250000 | 29.000000  | 3798.368899 | ... | 0          | 2023 | 3798.36 |
| 2463 | 2023-12-26 | 3780.100098 | 3833.850098 | 3780.100098 | 3794.600098 | 3794.600098 | 70216.0  | 53.750000 | 14.500000  | 3803.118699 | ... | 1          | 2023 | 3803.11 |
| 2464 | 2023-12-27 | 3795.550049 | 3818.000000 | 3768.100098 | 3810.800049 | 3810.800049 | 28290.0  | 49.899902 | 15.250000  | 3810.233399 | ... | 0          | 2023 | 3810.23 |
| 2465 | 2023-12-28 | 3822.100098 | 3838.250000 | 3793.750000 | 3801.050049 | 3801.050049 | 29256.0  | 44.500000 | -21.050049 | 3802.150065 | ... | 0          | 2023 | 3802.15 |

482 rows x 23 columns

```
In [73]: # Convert 'Date' column to datetime format
df_cleaned['Date'] = pd.to_datetime(df_cleaned['Date'])

# Plot the cumulative market and strategy returns
plt.figure(figsize=(18, 8))

plt.plot(df_cleaned['Date'], df_cleaned['Cumulative Market Returns'], color='r', label='Market Returns')
plt.plot(df_cleaned['Date'], df_cleaned['Cumulative Strategy Returns'], color='g', label='Strategy Returns')

plt.title('Cumulative Market and Strategy Returns')
plt.xlabel('Date')
plt.ylabel('Cumulative Returns')
plt.legend()
plt.grid(True)

plt.show()
```



GitHub Link (Python codes): <https://github.com/xxsarikapatel/AI-ML-1Coursework/blob/main/AI%26%20ML%20CW.ipynb>